FOR YOUR JOB SEARCH

Stop refreshing LinkedIn all day.

You already know the drill: refresh the search, scroll past a dozen postings that are obviously wrong, hope you don't miss the one that isn't. Do that every day for weeks, and it wears you down, and it's easy to miss something good in all the noise.

SAME SEARCH. LESS NOISE.

What it feels like now

Marketing Intern, unpaid, wrong stage entirely

Senior Backend Engineer, on-site only, wrong field

Staff Product Designer, actually a good fit

Executive Assistant, wrong seniority

Warehouse Associate, wrong industry

Customer Support Lead, below your salary range

Lead Product Designer, also worth a look

Sales Development Rep, wrong industry

What you get instead

#1 MATCH

Staff Product Designer

Remote, mid-market SaaS, right level, right salary band, no dealbreakers.

Strong fit

#2 MATCH

Senior Product Designer, Series B

Remote-first, matches your target range, one step outside preferred seniority.

Good fit

#3 MATCH

Lead Product Designer, Remote

Great scope and salary, but hybrid three days a week, worth a look.

Worth a look

HOW IT WORKS

Four steps, and none of them are you refreshing a page.

STEP 1

You describe what a good fit looks like

Seniority, location, salary, dealbreakers, the real criteria, not just a job title.

STEP 2

It watches for new postings, all day

Continuously, in the background, so you don't have to keep checking yourself.

STEP 3

Every posting gets scored against your criteria

Like having someone read each one and judge the fit, not just the keywords.

STEP 4

You see one ranked shortlist

Best matches first, instead of a raw feed you have to sift through yourself.

No more refreshing. Just a shortlist, waiting for you.

APPENDIX

Behind the scenes

A technical summary of the pipeline, for reference. This is the plumbing behind the four steps described earlier.

STAGE	WHAT HAPPENS
Watching LinkedIn	Two saved LinkedIn searches (specific title, seniority, and location criteria) are each turned into an RSS feed via rss.app, which polls LinkedIn's search results and detects new postings automatically.
Real-time capture	The moment a new posting is detected, a webhook pushes it to a small receiving endpoint in this project's webapp. The endpoint verifies the push and stores the raw posting (title, link, description, date) in a Notion database, which acts as the holding area, unscored.
Scoring, three times a day	An automated job reads everything unscored, skips duplicates, and fetches full posting details for anything plausibly relevant (to confirm remote/hybrid/onsite and pay). Each posting is scored 1–5 against a personal rubric covering seniority, industry fit, location, salary floor, and scope, with automatic disqualifiers for on-site-only roles outside the user's city or pay below the floor. The score and a one-sentence explanation are saved back to that Notion database, which serves as the results list. Postings older than a week are flagged stale; older than two weeks are archived.
Viewing results	The webapp, hosted on Vercel, reads that Notion database and displays it as a ranked, filterable list sorted best-fit-first, with a location filter. Clicking a listing opens a detail view with the full score breakdown and a link to the original posting.